

LLMs in Plain Language: The Storytelling Game

Day 2 morning · Playful Programming and Artificial Intelligence · University of Copenhagen

You are the model. Round by round, the room builds a sentence one word at a time, on the board or a large sheet everyone can see.

Materials: writing surface, one six-sided die, voting by hands.

Round 1 — Free story

5 min

“The last teacher on Earth opened the classroom door and...”

Go around the room. Each person adds exactly one word, any word they like. Usually surreal and funny — leave it on the board.

Round 2 — The probable story

10 min

New fragment

Each person proposes the most likely next word; the room votes between the top two proposals. Majority wins. The story becomes grammatical, coherent, and noticeably boring — that boredom is the point.

Round 3 — The temperature dial

10 min

New fragment

Same voting rule, but roll the die first. 1–4: take the majority word. 5: take the second-place word. 6: the current player may pick any word at all.

Round 4 — The confident lie

10 min

“The capital of Greenland is famous for its...”

Same probable-word rules. The room will produce something fluent about Nuuk that is partly or wholly invented — nobody in the game ever checked a fact.

Optional Round 5 — Meet the real one

5 min

Projector, chat interface

Give a chat model one of the same opening fragments three times and compare outputs. Then ask it something obscure enough to produce a hallucination.

TERMS NAMED AS THEY'RE EXPERIENCED

Sampling → temperature → hallucination

Storytelling Game — running notes

During the game

Keep the pace brisk. One word per person, no speeches — the game dies if word selection becomes a seminar.

Do not correct “wrong” votes. If the room picks an odd word in Round 2, the story bends and recovers — that's how models behave too.

The boredom of Round 2 is a feature. Do not apologise for it or cut it short; the contrast with Round 1 does the teaching.

Vocabulary discipline. Introduce temperature, sampling and hallucination only after the room has experienced the phenomenon.

Watch for anthropomorphism. “The model wants...”, “it's trying to...” — gently reformulate in prediction language each time.

If something goes wrong

Someone already knows the theory and narrates ahead. Recruit them as scribe or die-roller.

Round 4 produces something true by accident. Run one more fragment with a more obscure hook, e.g. the 1954 Copenhagen tram strike.

Group too large for one word per person. Split into two parallel stories and compare outputs at the end — demonstrates sampling variance.

What this simplifies

Real models predict tokens (word fragments), not whole words.

Real models weigh thousands of probability-ranked candidates, not two proposals.

The “voting” in a real model is a learned function of the preceding text, not a democratic show of hands.

Context windows, embeddings, and attention are entirely absent here.

Classroom transfer

Works from age ~13 upward with no changes; fragments can be tuned to any subject.

Language teachers can run Round 2 in the target language, doubling as grammar practice.

Plugged follow-up: the same experiment in a real interface with temperature visible, e.g. an API playground.
